

Telecommunications Assignment #3

Ciarán Moyne, Tom Murphy, Jonathan Rosca

January 10, 2014

0.1 Introduction

This is a system of distribution of files over a local network, using UDP multicast sockets for efficient one-to-many transmissions.

It does not require a server, but requires a router capable of multicast networking, so only operates on the local network. Every participant is called a peer. It is written in Python.

0.2 Protocol

0.2.1 Network membership

For a client to be part of a network of peers, it need only have access to other clients in the local computer network. Every client maintains a list of known peers. The network address is used to identify a peer, since it is unique within the local network.

In order to let the other peers know that a new client has joined the network, a small presence packet is sent to the multicast socket. When received by the other peers, they each add the new peer to their list, and keep track of when it was they received the last presence packet, as well as the nickname. The presence packet is sent periodically (every 200 ms) by each peer for their entire runtime. If a peer's presence was announced less recently than 5 seconds ago, they are thought to have left the network, and are deleted from the known peers list.

The packet also contains a nickname for the client, picked by the user. It cannot be longer than 30 unicode characters. It is only used to allow the users to identify peers in the network, and can be changed instantly. It does not have to be unique between the peers. It will be updated for the other peers when they receive the next presence packet from this client.

0.2.2 File distribution

Metadata

The relevant metadata of a file is the filename, size, type(image, video, audio or other), sha256 hash digest of the entire content and the time to live (picked by the user, not stored in the file). A file may optionally include a thumbnail of up to 90x90 32-bit RGBA pixels, where appropriate and possible.

Recipients

The recipients are the known peers at the time of the start of the transmission. If a new peer joins the network after the transmission has started, it is excluded from the transmission.

File chunks

The file contents are sent as chunks, which are as large as possible (almost 64KiB) in the UDP protocol. The system thus relies on IP packet fragmentation and error detection to drop incorrect packets. If the layout of a packet is correct, an error in the payload of an individual chunk packet cannot be detected.

Transmission

The metadata is sent to all recipients. When every recipient the metadata, every content chunk is sent until each one is acknowledged by every recipient.

Lifetime

The "life" of a distributed file on a peer's client starts when it was fully received. The time this happens can be different for each recipient, since different peers can have different chances of dropping packets. Each recipient must delete its copy of the file when the time is *arrived_at* + *tll*. For this reason, in case the distributing peer does not want to prematurely stop the distribution of the file, the receiving peers do not have to send deletion acknowledgement packets when they have deleted the file.

At any point, from the start of the transmission of metadata, the user might change their mind and wish to stop sending the file and have it deleted from the peers. The peer sends a deletion request to the recipients, repeatedly, until they each acknowledge having deleted their copy of the file, partial or complete.

0.3 Implementation

Transmission

There are three lists of current recipients, for metadata, chunks(content) and deletion acknowledgements, which contain which packets were not acknowledged yet(even if the packets weren't even sent). If a recipient times out, it is excluded from the transmission and deleted from the lists.

When all the chunks of a file are received, their hash is computed and compared to the one from the metadata. The latter is assumed to be correct, so if

the hashes do not match, the content is thought to be erroneous, and the file is discarded.

0.3.1 Drawbacks

If a file is completely received, and a single chunk is incorrect, causing the hash not to match the metadata hash, it cannot be salvaged and is dropped entirely. A possible solution would be to send a partial checksum along with every chunk packet, so that the recipient will know was the first erroneous chunk it received, and then send negative acknowledgements for every packet after that.

The program cannot deal with disconnecting from the network.